

Thread Creation: Parallel Regions

- You create threads in OpenMP* with the parallel construct.
- For example, to create a 4 thread Parallel region:

Each thread executes a copy of the code within the structured block

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID = omp_get_thread_num();
    pooh(ID,A);
}
```

Runtime function to request a certain number of threads

Runtime function returning a thread ID

- Each thread calls pooh(ID,A) for ID = 0 to 3

Thread creation: How many threads did you actually get?

- Request a number of threads with `omp_set_num_threads()`
- The number requested may not be the number you actually get.
 - An implementation may silently give you fewer threads than you requested.
 - Once a team of threads has launched, it will not be reduced.

Each thread executes a copy of the code within the structured block

```
double A[1000];
omp_set_num_threads(4); ← Runtime function to request a certain number of threads
#pragma omp parallel
{
    int ID      = omp_get_thread_num();
    int nthrds = omp_get_num_threads();
    pooh(ID,A);
}
```

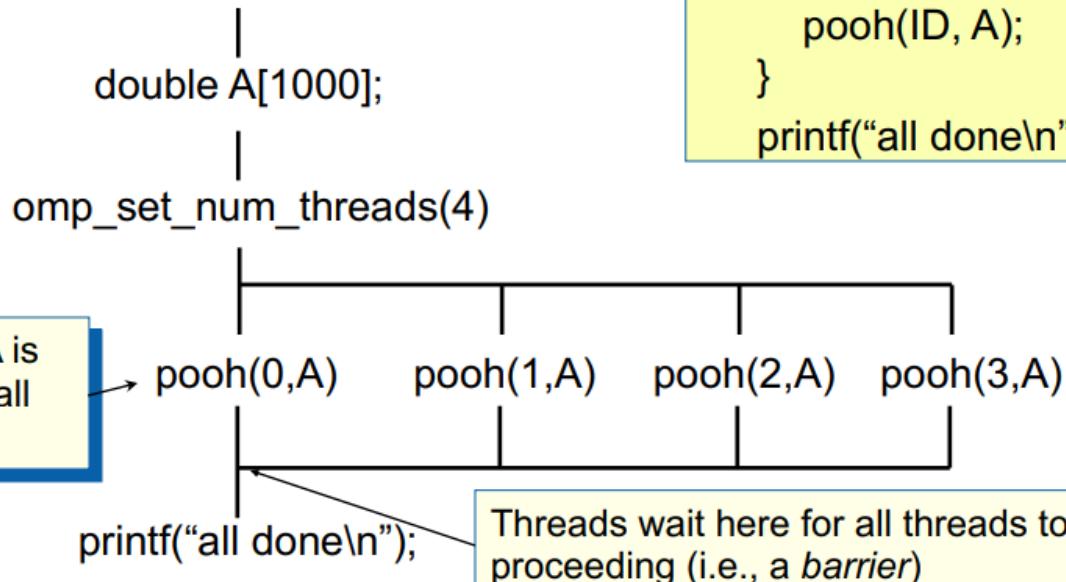
Runtime function to request a certain number of threads

Runtime function to return actual number of threads in the team

- Each thread calls `pooh(ID,A)` for $ID = 0$ to $nthrds-1$

Thread Creation: Parallel Regions Example

- Each thread executes the same code redundantly.



```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID = omp_get_thread_num();
    pooh(ID, A);
}
printf("all done\n");
```



OPENMP COMPONENTS

Compiler Directives

- A. **Parallel construct**
- B. **Work Sharing**
- C. **Synchronization**
- D. **Data Scope**
 - Private
 - Shared
 - First private
 - Last private
 - Reduction

Runtime Library

- **Number of threads**
- **Thread ID**

Environment Variables

- **Number of threads**
- **Scheduling Type**



OPENMP COMPILER DIRECTIVES

- Spawning a parallel region
- Used to divide blocks of code among threads.
- Distributing loop iterations among threads.
- Serializing sections of code.
- Synchronization of work among threads.



RUNTIME LIBRARY ROUTINES

- Setting and querying the number of threads.
- Querying unique thread ID.
- To check whether inside parallel region.

ENVIRONMENT VARIABLES

- Setting the number of threads.
 - Export OMP_NUM_THREADS=8
 - Specifying how loop iterations are divided.
 - Binding threads to processors



1. OPENMP COMPILER DIRECTIVES

- Syntax:

- `#pragma omp directive-name [clause1 clause2..] new-line`

- Example:

- `#pragma omp parallel`
- `#pragma omp parallel private(i,j)`

- Example Directives

- Parallel

- for/do

- Sections

- Single

- Critical

- Barrier

OpenMP Code Structure

- ▶ `#pragma omp <directive-name> [clause,clause,..] new-line\`
- ▶ General Rules
 - ▶ Case sensitive
 - ▶ Compiler Directives follow C/C++ standards
 - ▶ Only one directive name to be specified per directive

PARALLEL DIRECTIVE

- ▶ Purpose: creating a team of threads in the parallel region.
- ▶ Format:

```
#pragma omp parallel [clause...] newline  
    if (scalar_expression)  
    private(list)  
    shared(list)  
    firstprivate(list)  
    reduction(operator:list)  
    default(share | none)  
    copyin(list)  
    num_threads(integer-expression)
```



- Main thread **creates a team of threads** and becomes the master of the team.
 - The **master** is a member of that team and has **thread id 0** within that team.
 - Starting from the beginning of this parallel region, the **code is duplicated** and all threads will execute that code.
 - There is an **implied barrier** at the end of a parallel section.
 - If any thread terminates within a parallel region, all **threads** in the team **terminate**.
-
- Restrictions:
 - A parallel region must be a structured block that does not span multiple routines or code files
 - Only a **single IF clause** is permitted
 - Only a **single NUM_THREADS clause** is permitted



A. PARALLEL DIRECTIVE

- The parallel construct forms a team of threads and starts parallel execution of a parallel region.
- A parallel region is a block of code that will be executed by multiple threads.

```
#pragma omp parallel [clause1 clause2...] newline  
structured_block
```

Clauses:

if (scalar_expression)

private (list)

firstprivate (list)

num_threads (integer-expression)



C/C++ - General code Structure

```
#include <omp.h>
```

```
main () {
```

```
}     int var1, var2, var3;  
Serial code
```

```
.
```

```
.
```

```
.
```

Beginning of parallel section. Fork a team of threads.
Specify variable scoping

```
#pragma omp parallel private(var1, var2) shared(var3)
```

```
{
```

Parallel section executed by all threads

```
.
```

Other OpenMP directives

```
.
```

Run-time Library calls

```
.
```

All threads join master thread and disband

```
}
```

Resume serial code

```
.
```

```
}
```



Example: parallel construct

```
#include <omp.h>
main ()
{
int nthreads, tid;
/* Fork a team of threads with each thread having a private tid */
#pragma omp parallel private(tid)
{
    /* Obtain and print thread id */
    tid = omp_get_thread_num();
    printf("Hello World from thread = %d\n", tid);

    /* Only master thread does this */ if
    (tid == 0)
    {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
} /* Implicit Barrier - All threads join master thread and terminate */
}
```



- Main thread **creates a team of threads** and becomes the master of the team.
- The **master** is a member of that team and has **thread id 0** within that team.
- There is an **implied barrier** at the end of a parallel section.
- If any thread terminates within a parallel region, **all threads** in the team **terminate**.



DIFFERENT WAYS OF THREAD CREATION

- The number of threads in a parallel region is determined by the following factors, **in order of precedence (Low- High):**
 1. Default number of Threads (number of cores)
 2. Setting of the **OMP_NUM_THREADS** environment variable

```
export OMP_NUM_THREADS=2
```
 3. Use of the **omp_set_num_threads()** library function

```
omp_set_num_threads(4);
```
 4. Setting of the **NUM_THREADS** clause

```
#pragma omp parallel private(tid) num_threads(6)
```
- Implementation default - usually the **number of cores**.
- Threads are numbered from 0 (master thread) to N-1

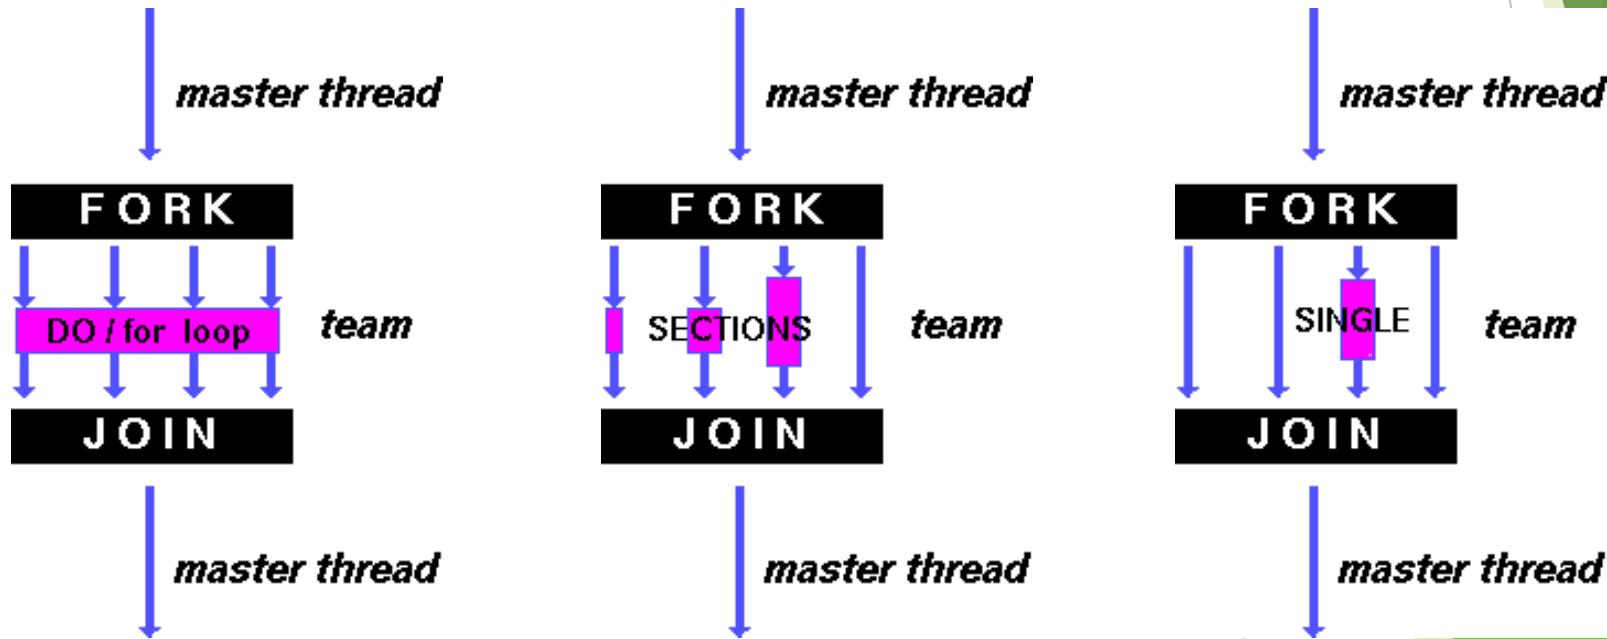


B. WORK SHARING CONSTRUCTS

- Divides **execution** of code region among members of the team of threads.
- Work-sharing constructs **do not launch new threads**
- Restrictions
 - Must be **enclosed** within a parallel region.
 - Work is **distributed** among the threads
 - Encountered by all threads
 - Does not launch new set of threads

Examples:

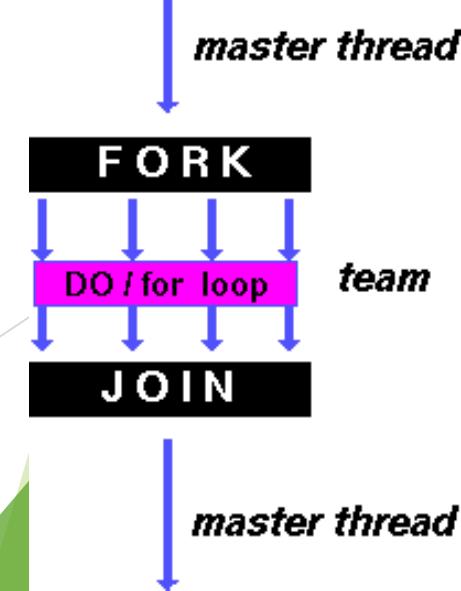
- For – data parallelism
- Section – functional parallelism
- single – serializes a section of code



FOR DIRECTIVE

- **for directive** divide the loop iterations among threads and execute in parallel.
- **Purpose:** **For directive** specifies that the iterations of the loop immediately following it must be executed in parallel by the team
- **Syntax:**

```
#pragma omp for [clause1      clause2    ...] newline
for_loop
```





Example: for construct

```
#include <omp.h>
#define N    1000

main ()
{
int i;
float a[N], b[N], c[N];

/* Some initializations */
for (i=0; i < N; i++)      a[i] = b[i] = i * 1.0;

#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp for for
    (i=0; i < N; i++)
        c[i] = a[i] + b[i];
} /* end of parallel section */
}
```



RESTRICTIONS

- `for (index = start ; index < end ; increment_expr)`
 - It should be possible to **determine** the number of loop **iterations** before execution, to divide the iterations.
- No while loops
- No variations of for loops where the start and end values change
- increment must be same each iteration
- It is illegal to branch(goto) out of a loop associated with a for directive



If

- A clause to decide at run time if a parallel region should actually be executed in parallel (multiple threads) or just by the master thread:

- Syntax:
`If(logical expr)`

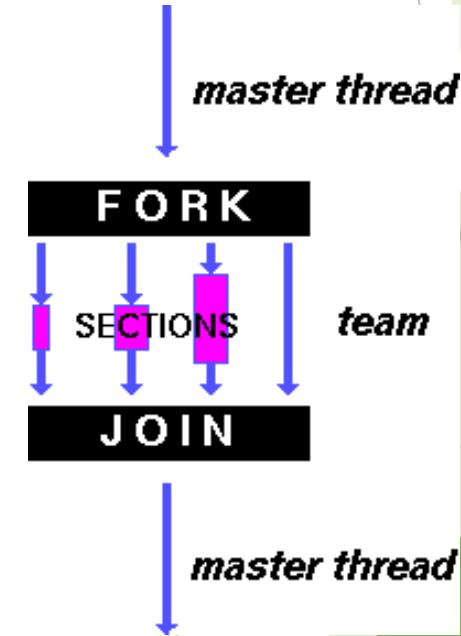
- Example:
`#pragma omp parallel if (n>100000)`



SECTION DIRECTIVE

- It specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- Each SECTION is executed only once by a thread in the team.
- Syntax:

```
#pragma omp sections [clause1 clause2 newline]
{
    #pragma omp section
        structured_block
    #pragma omp section <newline>
        structured_block
}
```





○ Key points

- NOWAIT : Implicit barrier at the end of sections directive, if nowait clause is present.
- If “too many” sections, some threads execute more than one section (round-robin).
- If “too few” sections, some threads are idle.
- We don’t know in advance which thread will execute which section.

○ Restrictions

- It is illegal to branch into or out of section blocks
- SECTION directives must occur within the lexical extent of an enclosing SECTIONS directive

```
#pragma omp parallel default(none)\\
shared(n,a,b,c,d) private(i)
{
    #pragma omp sections nowait
    {
        #pragma omp section
        for (i=0; i<n-1; i++)
            b[i] = (a[i] + a[i+1])/2;

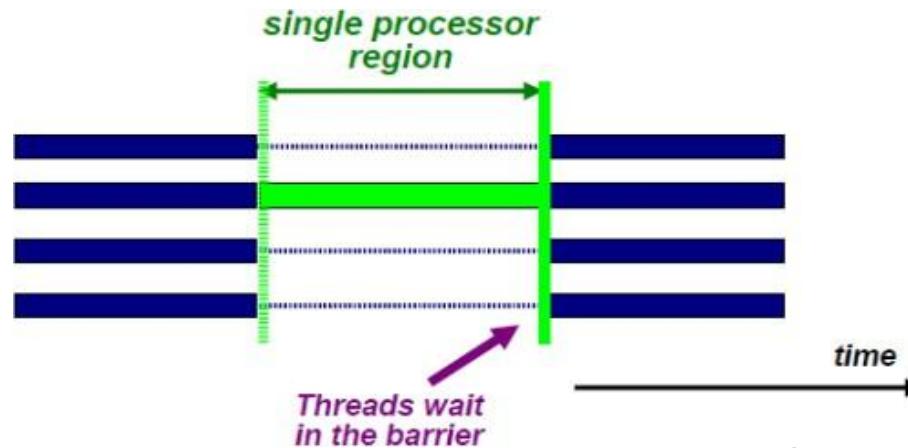
        #pragma omp section
        for (i=0; i<n; i++)
            d[i] = 1.0/c[i];
    } /*-- End of sections --*/
} /*-- End of parallel region --*/24
```

No need to wait
to enter Or to
exit from a
section

SINGLE DIRECTIVE

- The enclosed code is to be **executed by only one thread** in the team. May be useful when dealing with sections of code that are not thread safe such as I/O
- A **barrier** is implicitly set at the end of the single block (the barrier can be removed by the nowait clause)
- Syntax:

```
#pragma omp single [clause1 clause2 ...]  
structured_block
```

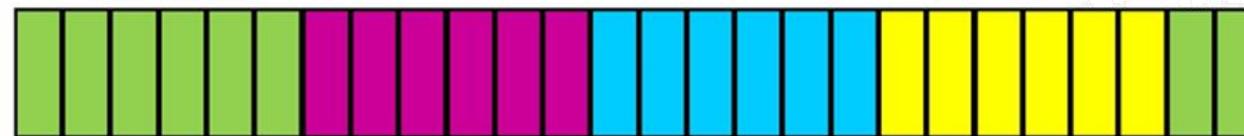




CLAUSES

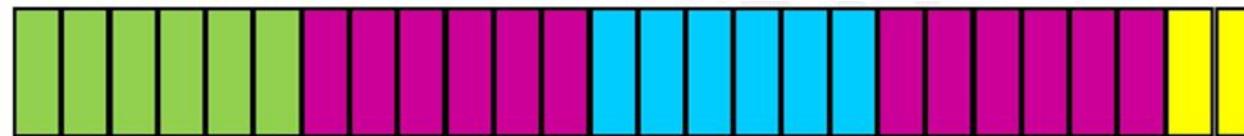
SCHEDULE: Describes how iterations of the loop are divided among the threads in the team. The default schedule is implementation dependent.

- **STATIC**



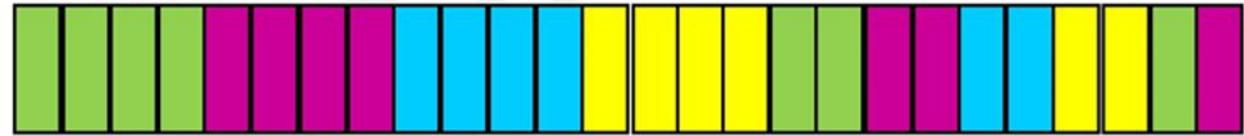
SCHEDULE(STATIC,6)
26 iter on 4 processors

- **DYNAMIC**



SCHEDULE(DYNAMIC,6)
26 iter on 4 processors

- **GUIDED**



SCHEDULE(GUIDED,4)
26 iter on 4 processors

- **RUNTIME** – determined by the environment variable **OMP_SCHEDULE**



SYNCHRONIZATION

- Synchronization is the task of ensuring that the multiple threads can safely share the resources.
- It is used to avoid the race condition

- Consider a case of 2 threads, both trying to increment a variable “x” at the same time (assume x = 0 initially)
- Problem: One possible execution sequence:
 - Thread 1 loads the value of x into register A.
 - Thread 2 loads the value of x into register A.
 - Thread 1 adds 1 to register A → $x = x + 1$ by thread1
 - Thread 2 adds 1 to register A → $x = x + 1$ by thread2
 - Thread 1 stores register A at location x
 - Thread 2 stores register A at location x
 - The resultant value of x will be 2, not 1 as it should be.



SYNCHRONIZATION

- Solution: To avoid a situation like this, the incrementing of x must be synchronized between the two threads to ensure that the correct result is produced.

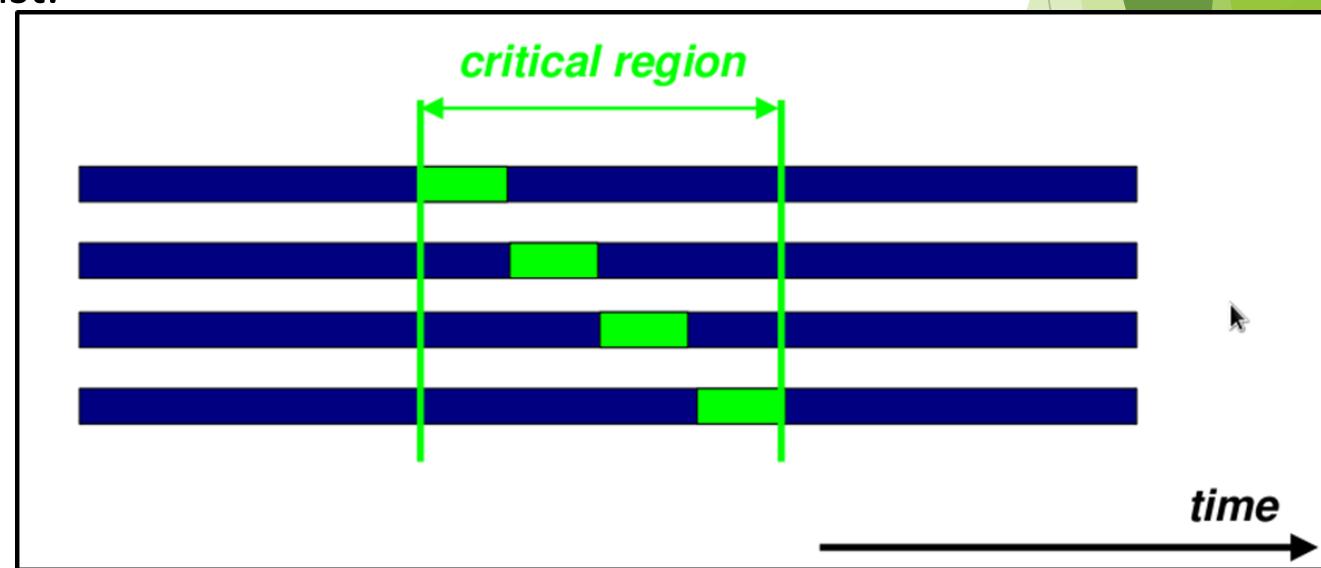
- OpenMP Synchronization Constructs:
 - CRITICAL Directive
 - ATOMIC Directive
 - BARRIER Directive
 - MASTER Directive

CRITICAL

- The CRITICAL directive specifies a region of code that must be **executed by only one thread at a time.**
- Critical section **prevents** multiple threads from accessing a section of code at the same time.
- Only **one active thread can** update the code inside the critical region.
- Syntax: `#pragma omp critical`

```
{ structured_block  
}
```

The optional name enables multiple different CRITICAL regions to exist.





CRITICAL

- If a thread is currently executing inside a CRITICAL region and another thread reaches that CRITICAL region and attempts to execute it, it will block until the first thread exits that CRITICAL region.
- The optional name enables multiple different CRITICAL regions to exist:
 - Names act as global identifiers. Different CRITICAL regions with the same name are treated as the same region.
 - All CRITICAL sections which are unnamed, are treated as the same section.



ATOMIC

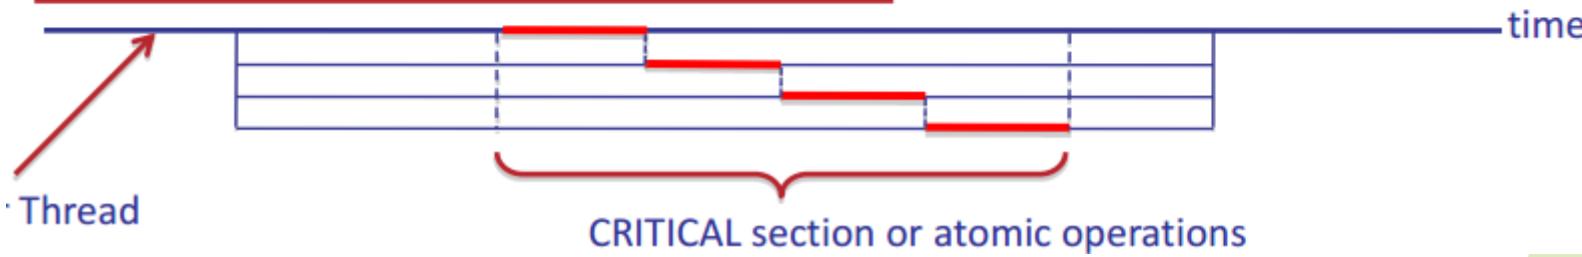
- Specifies that a specific memory location must be updated atomically, by a single thread
- It is used to prevent race conditions
- Atomicity: We cannot split an operation – Example, A write operation.
- The directive applies only to a single, immediately following statement → Applies to only a single statement, allows only a limited number of expressions.
- Provides a mini critical section
- Syntax:

```
#pragma omp atomic
{
    statement_expression
}
```

C.4 ATOMIC .VS. CRITICAL

```
#pragma omp parallel shared(sum,x,y)
...
#pragma omp critical
{
    update(x);
    update(y);
    sum=sum+1;
}
...
!$OMP END PARALLEL
```

```
#pragma omp parallel shared(sum)
...
{
#pragma omp atomic
    sum=sum+1;
}
...
```



- This Directive is very similar to the CRITICAL directive.
- Difference is that ATOMIC is only used for the update of a memory location.
- Sometimes ATOMIC is also referred to as a mini critical section.



BARRIER

- On reaching BARRIER directive, a **thread will wait** at that point until all other threads have reached that barrier.
- All threads then resume executing in parallel the code that follows the barrier. **Synchronizes all threads in a team.**
- Syntax:

```
#pragma omp barrier newline
```

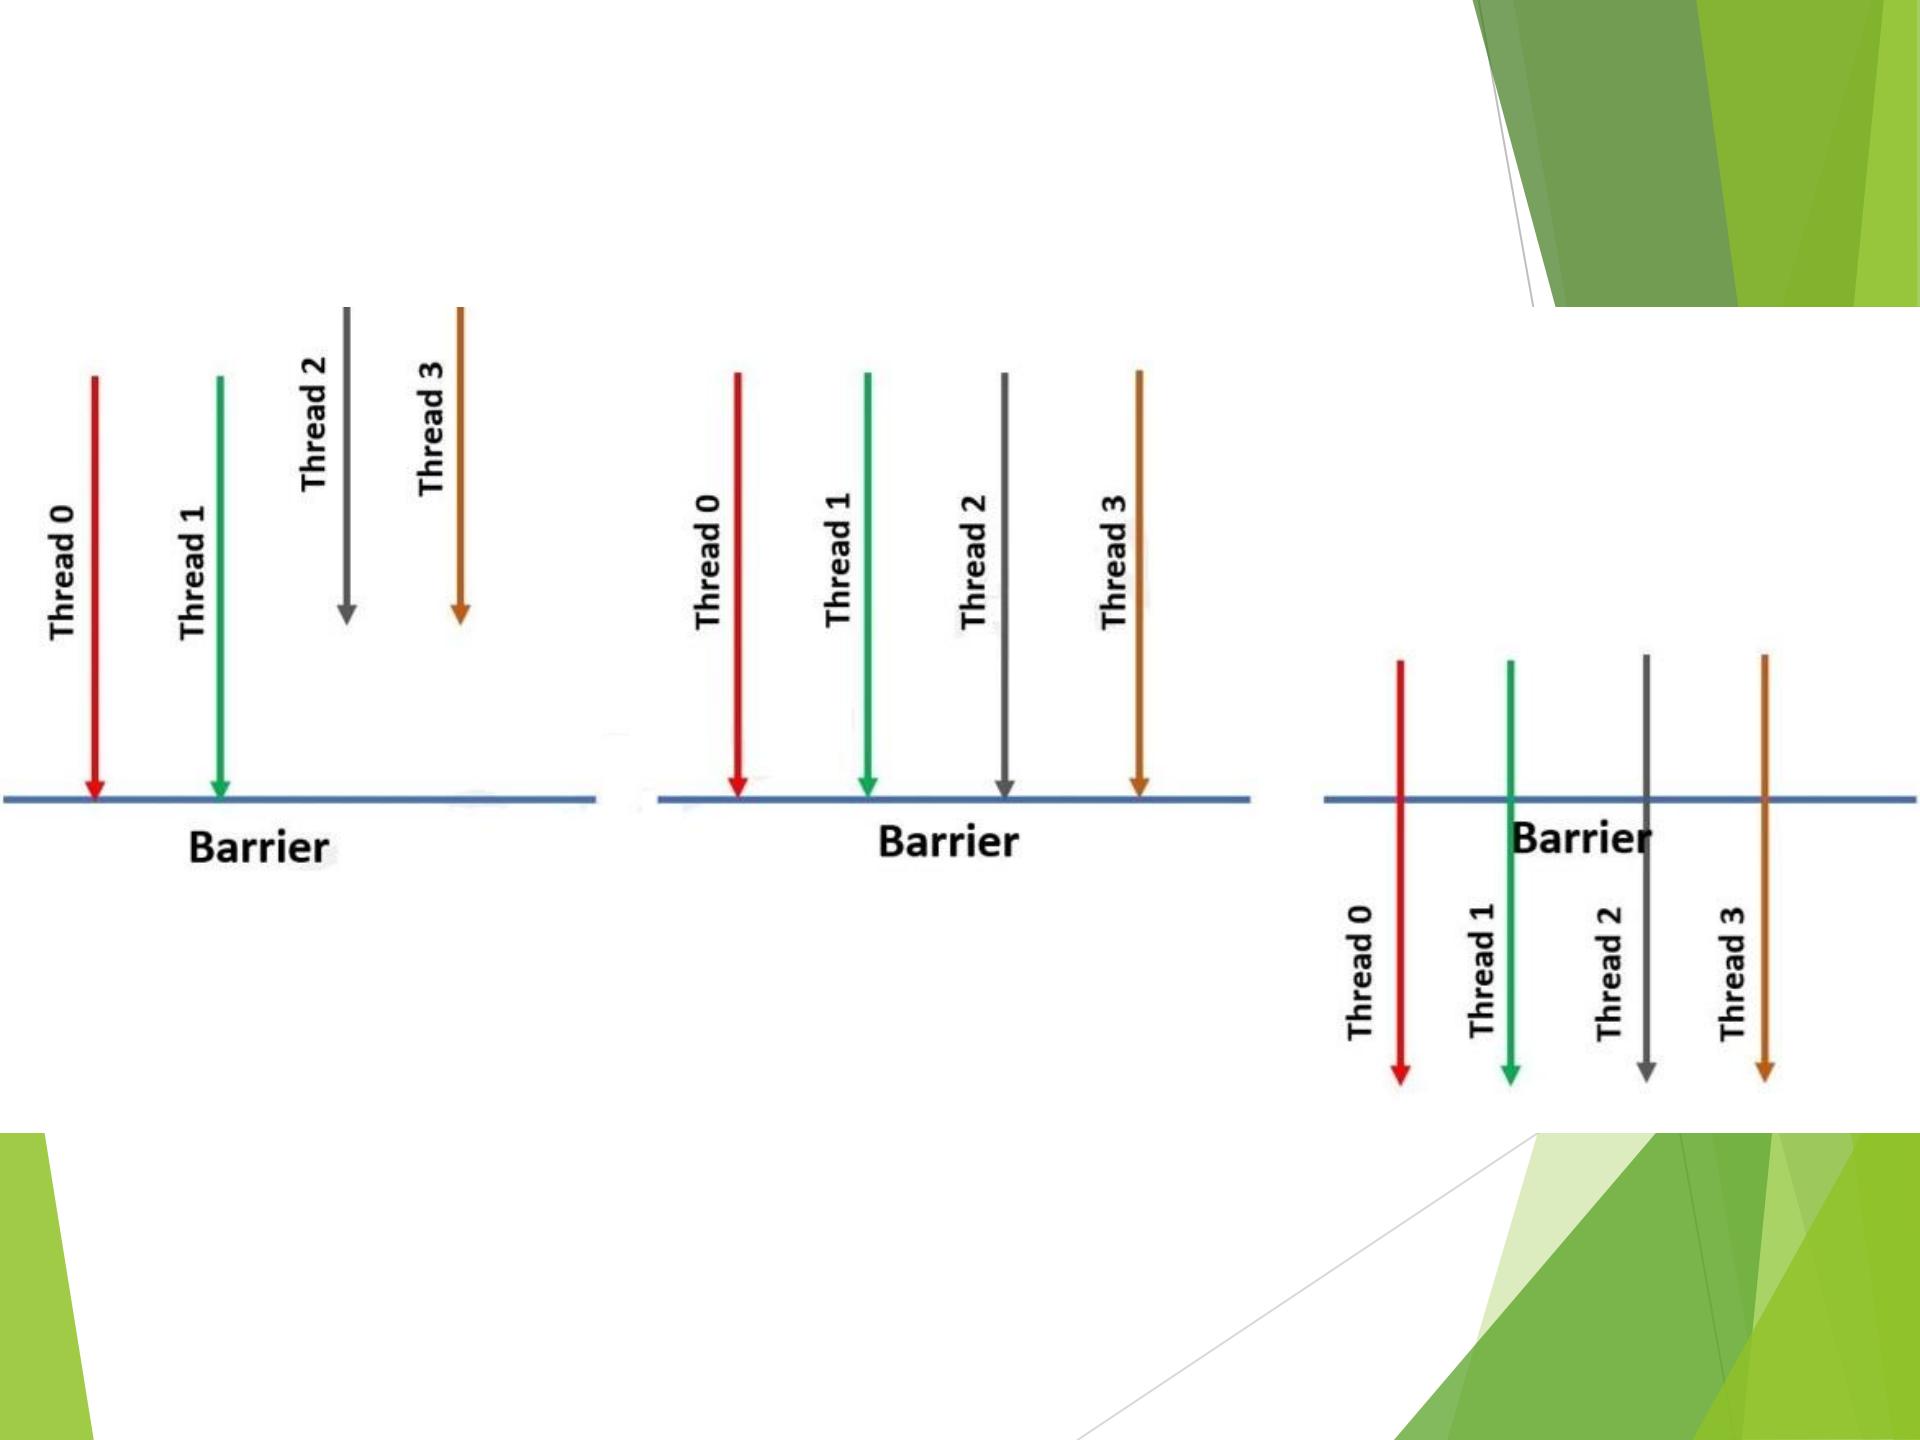
```
#pragma omp parallel for  
for (i=0; i < N; i++)
```

```
a[i] = b[i] + c[i];
```

```
#pragma omp barrier
```

```
#pragma omp parallel for  
for (i=0; i < M; i++)
```

```
d[i] = a[i] + b[i];
```





MASTER

- To execute a region only by master thread of the team.
- All other threads on the team skip this section of code
- There is no implied barrier associated with this directive
- Syntax:

```
#pragma omp master newline  
structured_block
```



DATA SCOPE ATTRIBUTES

The OpenMP Data Scope Attribute Clauses are used to explicitly define how variables should be scoped. They include:

- IF
- PRIVATE
- FIRSTPRIVATE
- SHARED

Data Scope Attribute Clauses are used in conjunction with several directives (PARALLEL, DO/for, and SECTIONS) to control the scoping of enclosed variables.



PRIVATE

- This declares variables in its list to be private to each thread
- Syntax:

`private (list)`

Eg: `int B = 10;`

`#pragma omp parallel private(B)`

`{`

`B=...;`

`}`

- A un-initialised copy of B is created before the parallel region begins.



FIRSTPRIVATE

- **Firstprivate:** A private initialized copy of B is created on each thread's stack before the parallel region begins

● Syntax:

firstprivate (list)

● Example:

```
int B = 10;
```

```
#pragma omp parallel firstprivate(B)
```

```
B = 10;
```



SHARED

- SHARED Clause

- A shared variable **exists in only one memory location** and all threads can read or write to that address

- Syntax:

shared (list)

- It is the programmer's responsibility to ensure that multiple threads properly access SHARED variables (such as via CRITICAL sections)



2. RUNTIME LIBRARIES

- Execution environment routines that can be used to control and to query the parallel execution environment

Example routines:

- OMP_SET_NUM_THREADS
 - omp_set_num_threads routine affects the number of threads to be used for subsequent parallel regions
 - C/C++ : void omp_set_num_threads(int num_threads);
- OMP_GET_NUM_THREADS
 - returns the number of threads in the current team.
 - C/C++ : int omp_get_num_threads(void);



- **OMP_GET_THREAD_NUM**

- Returns the thread ID of the thread

```
#include <omp.h>
int omp_get_thread_num()
```

- **OMP_GET_NUM_PROCS**

- To get the number of processors

```
#include <omp.h>
int omp_get_num_procs()
```

- **OMP_IN_PARALLEL**

- Determine if the section of code which is executing is parallel or not.

```
#include <omp.h>
int omp_in_parallel()
```



3. ENVIRONMENT VARIABLES

- OMP_SCHEDULE
 - export OMP_SCHEDULE "static"
 - export OMP_SCHEDULE "dynamic"

- OMP_NUM_THREADS
 - export OMP_NUM_THREADS=8

Data Scoping Attribute Clauses

- ▶ Purpose : The OpenMP Data Scope attribute clauses are used to explicitly define how variables should be scoped. They include:
 - ▶ Shared
 - ▶ PRIVATE
 - ▶ DEFAULT
 - ▶ FIRSTPRIVATE
 - ▶ LASTPRIVATE
 - ▶ COPYPRIVATE
 - ▶ COPYIN
 - ▶ REDUCTION

Shared and Private clauses

- ▶ Shared: There exists one instance of this variable which is shared among all the threads
- ▶ Private: Each thread in a team of threads has its own local copy of the private variable.
- ▶ Syntax: `shared(list)`
`private(list)`

Implicit Rules

```
int i=0;  
int n=10;  
int a=7;  
#pragma omp parallel for  
For(i=0;i<n;i++)  
{  
    int b=a+i;      // Here b is  
private variable and assigns its local copy to  
each thread  
}
```

- ▶ i, n, a and b are 4 variables
- ▶ n, a are shared variables
- ▶ Loop iteration variables are private by default. So i is private

- ▶ Variables which are declared locally within the parallel region are private. Thus b is private

```
int n=10;           //shared
int a=7;           //shared
#pragma omp parallel for
For(int i=0;i<n;i++) //i private
{
    int b=a+i;     // b private
}
```

- ▶ It is better to declare the loop iteration variables inside the parallel region

Explicit Rules

Shared

```
#pragma omp parallel for shared(n,a)
for(int i=0; i<n; i++)
{
    int b=a+i;      //
}
```

- ▶ The **shared(list)** clause declares that all variables in **list** are **shared**
- ▶ Shared variables introduce the overhead coz an instance of a variable is shared multiple threads
- ▶ Minimize the number of shared variables

Explicit Rules

Private

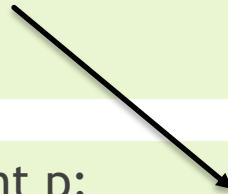
```
#pragma omp parallel for shared(n, a) private(b)  
for(int i=0; i<n; i++)  
{  
    int b=a+i;      //b is private  
}
```

- ▶ The **private(list)** clause declares that all variables in **list** are **private**
- ▶ OpenMP replicates the private variable and assigns its local copy to each thread.
- ▶ Here n, a are shared variables
- ▶ And b is private variable

Explicit Rules

Private

```
int p;  
  
#pragma omp parallel private(p)  
  
for(int i=0; i<n; i++)  
  
{  
  
    p=omp_get_thread_num();  
  
}
```



- ▶ We can avoid listing private variables inside the construct inside the parallel region
- ▶ Better to declare inside the parallel region

```
int p;  
  
#pragma omp parallel  
  
for(int i=0; i<n; i++)  
  
{  
  
    int p = omp_get_thread_num();  
  
}
```

Default

Default(shared)

```
int a, b, c, n;  
#pragma omp parallel for default(shared)  
for(int i=0; i<n; i++)  
{  
    .....  
}
```

- ▶ Observe two uses of Default clause
- ▶ Here, a, b, c, n are shared

```
int a, b, c, n;  
#pragma omp parallel for default(shared) private(a, b)  
for(int i=0; i<n; i++)  
{  
    // a and b are private variables  
    // c and n are shared variables  
}
```

- ▶ Here, is another use of default(shared)
- ▶ Here, c and n are shared
- ▶ And, a,b are private

Default Default(none)

```
int a, b, c, n;  
#pragma omp parallel for default(none) shared(b,c ,n)  
for(int i=0; i<n; i++)  
{  
    ....  
}
```

- ▶ None clauses forces to explicitly specify the data sharing attributes of all variables
- ▶ Error as var a data attribute is not specified

```
int a, b, c, n;  
#pragma omp parallel for default(none) shared(a, b, c, n)  
for(int i=0; i<n; i++)  
{  
    // a and b are private variables  
    // c and n are shared variables  
}
```

- ▶ Here, var a is included in shared
- ▶ Always write parallel regions with the default(none) clause
- ▶ Declare private variables inside parallel regions whenever possible

Critical Section

- ▶ Critical section prevents multiple threads from accessing a section of code at the same time.
- ▶ Only one active thread can update the data referenced by the code.
- ▶ SYNTAX:

```
#pragma omp critical [(name)]  
{  
    code_block  
}
```

- ▶ If critical sections are not used then manually acquire and release the locks
- ▶ What could be an example of critical section
 - ▶ Largest number from an Array of 10 numbers

Critical Section

- ▶ Critical section prevents multiple threads from accessing a section of code at the same time.

```
#include<omp.h>
int count=0;
Void myFunction()
{
    #pragma omp critical
    count++;
}
```

- ▶ What could be an example of critical section
 - ▶ Largest number from an Array of 10 numbers

Atomic

- ▶ Accessing specific memory location
- ▶ Race condition can be avoided
- ▶ SYNTAX:

```
#pragma omp atomic update  
{  
    expression statement;  
}
```

FirstPrivate clause

Purpose: FirstPrivate will copy the variable's original value in each thread's private variable

Syntax: firstprivate(var1, var2)

Here, var1 , var2 have instances in each thread.

LastPrivate clause

Purpose: Value from the last loop iteration upon exit from the parallel block is assigned to the original variable

Syntax: `lastprivate(var1, var2)`

Here, `var1` , `var2` have instances in each thread.

ThreadPrivate clause

Purpose: It makes the variables allocated to the threads persistent across the parallel regions

It also retains the computed values

It preserves the local copy for each thread

Syntax: `threadprivate(var1, var2)`

Here, var1 , var2 have instances in each thread.

Copyin clause

Purpose: It copies the values from the master thread to all other threads

Its only applicable for threadprivate variables

Syntax: copyin(var1, var2)

Here, var1 , var2 have instances in each thread.

Copyprivate clause

Purpose: The COPYPRIVATE clause can be used to broadcast values acquired by a single thread directly to all instances of the private variables in the other threads.

Associated with the SINGLE directive

Syntax: copyprivate(var1, var2)

Here, var1 , var2 have instances in each thread.

Conditional If in OpenMP

Purpose: It tells if the code inside the parallel region is executed in parallel or sequential

Syntax: `if(expression)`

if expression is true, the code inside parallel region
executed in parallel else in sequential.

Reduction clause:

Reduction is the accumulation of partial results from each threads into a single final result

A Private copy for each list variables is created for each thread

Syntax: reduction(operator : list)

Operators used with reduction clause:

Symbol	Meaning
+	Summation
-	Subtraction
*	Product
&	Bitwise AND
	Bitwise OR
^	shift
&&	Logical AND
	Logical OR

Clause / Directive summary:

Clause	Directive					
	PARALLEL	DO/for	SECTIONS	SINGLE	PARALLEL DO/for	PARALLEL SECTIONS
IF	◆				◆	◆
PRIVATE	◆	◆	◆	◆	◆	◆
SHARED	◆	◆			◆	◆
DEFAULT	◆				◆	◆
FIRSTPRIVATE	◆	◆	◆	◆	◆	◆
LASTPRIVATE		◆	◆		◆	◆
REDUCTION	◆	◆	◆		◆	◆
COPYIN	◆				◆	◆
COPYPRIVATE				◆		
SCHEDULE		◆			◆	
ORDERED		◆			◆	
NOWAIT		◆	◆	◆		