

# **Teaching Guidelines for**

## **Advanced Programming in HPC**

### **ACC-HPC June 2025**

**Duration:** classroom hours 36 + 36 lab hours

**Objective:** To introduce the student to Accelerator Programming OpenACC, OpenCL/SYCL ,CUDA Programming and Tools.

**Prerequisites:** Knowledge of programming in any language like C, C++, basic Mathematics, statistics.

**Evaluation method:** CCEE exam – 40% weightage

Lab exam – 40% weightage

Internal exam – 20% weightage

#### **List of Books / Other training material**

- Reference Book:**
1. OpenACC for Programmers: Concepts and Strategies by Sunita Chandrasekaran, Guido Juckeland
  2. Programming in Parallel with CUDA: A Practical Guide by Richard Ansorge
  3. Hands-On Gpu Programming with Python and Cuda: Explore high-performance parallel computing with CUDA by Dr Brian Tuomanen Dr. Brian Tuomanen (Author) Python Power: The Comprehensive Guide
  4. Learn CUDA Programming: A beginner's guide to GPU programming and parallel computing with CUDA 10.x and C/C++ by Jaegeun Han, Bharatkumar Sharma

Tools for High Performance Computing: Proceedings of the 2nd International Workshop on Parallel Tools for High Performance Computing, July 2008, HLRS, Stuttgart by Rainer Keller, Valentin Himmler, Bettina Krammer , Alexander Schulz

*Note: Each session mentioned is for theory and of 2 hours duration.*

*Lab assignments are indicatives, faculty need to assign more assignments for better practice.*

#### **Introduction to Accelerator Programming**

##### **Session 1:**

- Introduction GPU & GPGPU

##### **Session 2:**

- Performance of GPU over CPU

### **Session 3:**

- GPGPU hardware architecture (NVIDIA – AMD)
- Application of Accelerator Programming

### **Lab Assignment**

## **OpenACC**

### **Session 4:**

- Introduction to OpenACC
- Parallelize & Optimize Loops

### **Lab Assignment**

## **Session 5**

- Directive & Clauses
- GPUs and Multi-Core CPUs

### **Lab Assignment**

## **OpenCL/SYCL**

### **Session 6,7,8 & 9**

- Introduction to OpenCL/SYCL, Host Program & Device Program, Single-source programming for both host and device code, Queues to manage execution and synchronization, Buffers and accessors for data management between the host and device, Kernel programming, Parallelism via constructs, Platform portability across CPUs, GPUs, FPGAs, and other accelerators.

**Lab:** Case-Studies (Algorithms and Parallelization Approaches): Matrix – Multiplication

## **CUDA**

### **Session 10**

- NVIDIA GPGPU Architecture

### **Lab Assignment**

### **Session 11 & 12**

- GPGPU computing with CUDA

### **Lab Assignment**

## **Session 13 & 14**

- Thread & Memory Hierarchy
- CUDA constructs

### **Lab Assignment**

#### **Tools**

#### **Session 15:**

- Profiling: GNU (gprof), Intel VTune
- Debugger: GDB (GNU)

### **Lab Assignment**

#### **Session 16:**

- Performance measurement and optimization
- Algorithm, memory & network, parallel Processing, compiler flags, libraries (BLAS, LAPACK, MKL)

### **Lab Assignment**

#### **Session 17:**

- Benchmarking
- Introduction to micro-architecture analysis tools: LIKWID

### **Lab Assignment**

#### **Session 18:**

- Intel Advisor
- Demo of the sample code by using the above tools

### **Lab Assignment**